

Vinay Gundamaneni

SENIOR DATA ENGINEER

CONTACT

Vinay.g0228@gmail.com
952-681-0361
linkedin.com/in/
vinay-gundamaneni

CORE SKILLS

Languages

Python · SQL · PySpark · Bash

Data Stack

dbt · Apache Spark · Delta Lake
· Snowflake · Iceberg ·
BigQuery

Orchestration

Airflow · GCP Cloud Composer ·
Azure Data Factory

Cloud

GCP (BigQuery, Dataflow, Pub/
Sub, Composer) · Azure
(Databricks, Synapse, ADLS, ADF)

Streaming

Spark Streaming · GCP Dataflow
· Kafka · Pub/Sub

Reliability

Data contracts · SLOs · Lineage
tracking · Observability · CI/CD

Modeling

Medallion architecture · SCD Type
2 · Dimensional modeling

EDUCATION

MBA — Information Technology Services & Management

Concordia University, MN

PROFESSIONAL SUMMARY

Senior Data Engineer with 12+ years building production-critical data infrastructure at enterprise scale — specializing in billing, cost-center, and usage analytics pipelines that drive executive reporting and operational self-serve. Expert in Python, dbt, and Apache Spark across both GCP and Azure, with a cloud-agnostic approach that prioritizes reliability, full observability, and end-to-end ownership. Known for setting engineering standards, mentoring teams, and translating complex reporting requirements into scalable, well-governed data systems.

EXPERIENCE

Senior Data Engineer

Sep 2022 – Present | Remote

Total Quality Logistics (TQL)

- Own end-to-end pipelines for revenue, billing, and cost-center analytics — delivering gold-layer datasets that Operations and executive teams use daily for trusted self-serve reporting.
- Designed Medallion architecture (Bronze → Silver → Gold) in Delta Lake with enforced immutability — enabling point-in-time queries, full lineage, and audit-trail compliance across all production datasets.
- Built dbt transformation layer on Snowflake for cost-center and revenue reporting: SCD Type 2 tracking, referential integrity tests, custom quality macros, and freshness monitoring.
- Engineered Airflow DAG orchestration to replace brittle cron scheduling — introduced observable, retry-capable workflows with dependency management and full execution history.
- Defined data engineering standards for the org: CI/CD via Azure DevOps, dbt Cloud CI, and automated quality gates that block breaking changes before they reach any downstream consumer.
- Spearheaded org-wide migration to Apache Spark on Azure Databricks — built the proof-of-concept, secured leadership buy-in, and rolled out to production as the team's new engineering baseline.
- Mentored engineers on modeling, pipeline reliability, and observability — raising technical quality across all analytical workloads.

Python · PySpark · Spark · Snowflake · dbt · Airflow · Delta Lake · Azure Databricks · Azure DevOps · SQL

Data Engineering Lead

Apr 2021 – Sep 2022 | Remote

Prime Healthcare

- Designed and owned the data auditability framework for billing, accounts receivable, and cost-per-department analytics — serving Accounting and executive stakeholders under HIPAA-governed controls with full row-level traceability.
- Built dbt models with SCD Type 2 for billing history and cost hierarchies — enabling period-over-period variance analysis with contract validation enforced before any dataset reached a reporting consumer.
- Led a cross-departmental initiative across billing, clinical cost, and operations — owned data contract definitions, delivery milestones, and alignment across four business units.

- Migrated legacy SQL Server/SSIS pipelines to Apache Spark on Azure Databricks, shifting from overnight batch to near real-time processing and enabling intraday visibility for the first time.
- Built a self-serve analytics layer on GCP BigQuery and Azure Synapse with governed access controls — allowing Accounting and Operations to query trusted data independently.

Python · Spark · PySpark · dbt · Azure Databricks · Delta Lake · Airflow · GCP BigQuery · Azure Synapse · SSIS · SQL Server

Senior Data Engineer — Supply Chain & Analytics

Apr 2019 – Apr 2021 | Minneapolis, MN

Cognizant Technology Solutions — Client: Cargill Inc.

- Built pipelines on GCP (BigQuery, Dataflow, Pub/Sub) and Azure integrating Oracle ERP, Hadoop, and SQL Server into a unified analytics platform — serving Procurement and Operations globally with cost-center, purchase order, and vendor payment data.
- Designed a dimensional data warehouse with SCD Type 2 for vendor hierarchies, cost center histories, and PO lifecycle — enabling accurate period-over-period comparisons across global markets.
- Built Airflow DAG orchestration for multi-step workflows spanning global source systems, with dependency management, retry logic, alerting, and full execution audit logs.
- Mentored engineering teams across NA and APAC on modeling and pipeline reliability — driving consistent data quality across Oracle, Hadoop, and GCP-sourced datasets.

GCP BigQuery · GCP Dataflow · Pub/Sub · Python · Spark · Airflow · dbt · Oracle ERP · Hadoop · Azure Data Factory · SCD Type 2

Data Engineer

Mar 2016 – Dec 2017 | Hyderabad, India

Primera Medical Technologies

- Built ELT pipelines for EHR data extraction (billing, revenue cycle) from Meditech and Epic Clarity, feeding auditable reporting systems with HIPAA-compliant access controls.
- Designed a self-serve Power BI layer with row-level security, enabling clinical and operations teams to access trusted datasets independently.

SQL Server · Python · SSIS · Power BI · Meditech · Epic Clarity

Business Intelligence & Data Engineer

Aug 2012 – Feb 2016 | Hyderabad, India

Porvus Tech Solutions

- Developed SSIS ETL pipelines, SSRS reports, and SSAS dimensional models for healthcare and enterprise clients — data warehouse design, performance tuning, and stakeholder BI delivery.

SQL Server · SSIS · SSRS · SSAS · dimensional modeling